

Deletions from Red-Black Trees

Michael T. Gastner (28 March 2023)



Disclaimer: These slides are based on and occasionally quote from 'Introduction to Algorithms' (3rd ed.) by Cormen et al. (2009), MIT Press.

Deleting a Node z from a General Binary Search Tree

Reminder


Before we study the special case of red-black trees, let us review how to delete a node z from a general binary search tree. There are four different cases:

1. z has no left child.
2. z has a left child but not a right child.
3. z has two children and the successor of z (i.e., the node with the next larger key) is z 's right child.
4. z has two children and the successor of z is not z 's right child.

The next 4 pages illustrate how to delete a node in each of these cases.

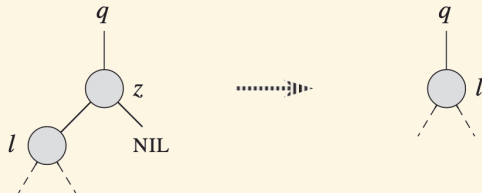
Deleting a Node z from a General Binary Search TreeCase 1: z Has No Left Child

If z has no left child, then we replace z with its right child r , which may or may not be NIL.



Deleting a Node z from a General Binary Search TreeCase 2: z Has a Left Child but No Right Child

If z has a left child l but no right child, then we replace z with l .




Deleting a Node z from a General Binary Search Tree

Case 3: z Has Two Children and the Successor of z Is z 's Right Child

If z has two children and its left child l is its successor y , then

1. we replace z with y .
2. we update y 's left child to be l .
3. y keeps its right child x .




Deleting a Node z from a General Binary Search Tree

Case 4: z Has Two Children and the Successor of z Is Not z 's Right Child

If z has a left child l and a right child r and if z 's successor y is not r , then

1. we replace y with y 's right child x .
2. we set y to be r 's parent.
3. we set y 's parent to be z 's former parent q , which may or may not be NIL.
4. we set l 's parent to be y .



Deleting a Node z from a General Binary Search Tree

TREE-DELETE, TRANSPLANT, and TREE-MINIMUM

In the pseudocode on the next page, the procedure $\text{TREE-DELETE}(T, z)$ deletes a node z from a binary search tree T . TREE-DELETE relies on two other procedures:

- $\text{TRANSPLANT}(T, u, v)$ replaces the subtree rooted at node u with the subtree rooted at node v .
- $\text{TREE-MINIMUM}(x)$ returns a pointer to the minimum element in the subtree rooted at node x .

The pseudocode of TRANSPLANT and TREE-MINIMUM is included on the next page alongside the pseudocode of TREE-DELETE .

Deleting a Node z from a General Binary Search Tree

Pseudocode

TREE-DELETE(T, z)

```
1  if  $z.left == \text{NIL}$ 
2      TRANSPLANT( $T, z, z.right$ )
3  elseif  $z.right == \text{NIL}$ 
4      TRANSPLANT( $T, z, z.left$ )
5  else  $y = \text{TREE-MINIMUM}(z.right)$ 
6      if  $y.p \neq z$ 
7          TRANSPLANT( $T, y, y.right$ )
8           $y.right = z.right$ 
9           $y.right.p = y$ 
10     TRANSPLANT( $T, z, y$ )
11      $y.left = z.left$ 
12      $y.left.p = y$ 
```

TRANSPLANT(T, u, v)

```
1  if  $u.p == \text{NIL}$ 
2       $T.root = v$ 
3  elseif  $u == u.p.left$ 
4       $u.p.left = v$ 
5  else  $u.p.right = v$ 
6  if  $v \neq \text{NIL}$ 
7       $v.p = u.p$ 
```

TREE-MINIMUM(x)

```
1  while  $x.right \neq \text{NIL}$ 
2       $x = x.right$ 
3  return  $x$ 
```


Running Time of a Deletion from a General Binary Search Tree

 $O(h)$

Each line of `TREE-DELETE`, including the calls to `TRANSPLANT`, takes constant time, except for the call to `TREE-MINIMUM` in line 5.

`TREE-MINIMUM` need $O(h)$ time on a tree of height h because the sequence of nodes encountered forms a simple path downward from the root.

Deleting a Node z from a Red-Black Tree


Keeping Track of Moving Nodes

When we delete a node z from a red-black tree, we must ensure that the resulting tree satisfies all five red-black properties. There are two important nodes, y and x that play an important role for maintaining the red-black properties:


- Node y :
 - ▶ If z has fewer than two children, then we set $y = z$.
 - ▶ If z has two children, then y is z 's successor, which is the node that moves into the position vacated by z .
- Node x is the node that moves into y 's original position.

We Cannot Directly Apply TREE-DELETE to Red-Black Trees


Scenarios that Violate Red-Black Properties



The root may become red.



Red parents may have red children.



Simple downward paths to different leaf nodes may contain a different number of black nodes.

Changes to TREE-DELETE

New procedures for repairing red-black properties and transplanting subtrees

To prevent violations of red-black properties, the procedure RB-DELETE will call a procedure named RB-DELETE-FIXUP.

RB-DELETE will also call a procedure RB-TRANSPLANT which differs in detail from TRANSPLANT. Here are the pseudocodes side by side with changes highlighted:

RB-TRANSPLANT(T, u, v)

```
1  if  $u.p == \textcolor{red}{T.nil}$ 
2       $T.root = v$ 
3  elseif  $u == u.p.left$ 
4       $u.p.left = v$ 
5  else  $u.p.right = v$ 
6       $v.p = u.p$ 
```

TRANSPLANT(T, u, v)

```
1  if  $u.p == \textcolor{red}{NIL}$ 
2       $T.root = v$ 
3  elseif  $u == u.p.left$ 
4       $u.p.left = v$ 
5  else  $u.p.right = v$ 
6  if  $v \neq \textcolor{blue}{NIL}$ 
7       $v.p = u.p$ 
```

RB-DELETE

Pseudocode

RB-DELETE(T, z)

```
1   $y = z$ 
2   $y\text{-original-color} = y.\text{color}$ 
3  if  $z.\text{left} == T.\text{nil}$ 
4       $x = z.\text{right}$ 
5      RB-TRANSPLANT( $T, z, z.\text{right}$ )
6  elseif  $z.\text{right} == T.\text{nil}$ 
7       $x = z.\text{left}$ 
8      RB-TRANSPLANT( $T, z, z.\text{left}$ )
9  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10      $y\text{-original-color} = y.\text{color}$ 
11      $x = y.\text{right}$ 
12     if  $y.p == z$ 
13          $x.p = y$ 
14     else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15          $y.\text{right} = z.\text{right}$ 
16          $y.\text{right}.p = y$ 
17     RB-TRANSPLANT( $T, z, y$ )
18      $y.\text{left} = z.\text{left}$ 
19      $y.\text{left}.p = y$ 
20      $y.\text{color} = z.\text{color}$ 
21     if  $y\text{-original-color} == \text{BLACK}$ 
22         RB-DELETE-FIXUP( $T, x$ )
```

Comment on RB-DELETE

Why Do We Not Need to Fix Up the Tree If *y-original-color* Is Red?

If *y-original-color* is red, the red-black properties still hold after line 20, for the following reasons:

- If y is red, then removing or moving it does not change any black-height in the tree.
- If y is red, then removing $y = z$ in the **if** or **elseif** clause cannot cause red nodes to become adjacent. If, instead, y is moved to z 's position in the **else** clause starting on line 9, then y takes z 's previous color on line 20, and, again, the action taken on y does not cause a red parent to have a red child.
- If y was not z 's right child, then y 's original right child x replaces y in the tree. If y is red, then x must be black; hence, replacing y by x cannot cause two red nodes to become adjacent.
- Because y could not have been the root if it was red, the root remains black.

How Do We Fix Up the Tree If *y-original-color* Is Black?

Case Distinctions

If x is red, we simply turn it black. This action permits x to be the root or have a red parent. Because the black node y is replaced by a black node x , no black-heights are changed.

If x is black, we must apply different actions depending on the colours of x and its relatives:

- Case 1: x 's sibling w is red.
- Case 2: x 's sibling w is black, and both of w 's children are black.
- Case 3: x 's sibling w is black, w 's left child is red, and w 's right child is black.
- Case 4: x 's sibling w is black, and w 's right child is red.

The actions may need to be taken iteratively. By the end, we ensure that x is black. The pseudocode is on the next two pages. Afterwards, I illustrate each case.

RB-DELETE-FIXUP

Pseudocode (Part 1)

RB-DELETE-FIXUP(T, x)

```
1  while  $x \neq T.root$  and  $x.color == BLACK$ 
2      if  $x == x.p.left$ 
3           $w = x.p.right$ 
4          if  $w.color == RED$ 
5               $w.color = BLACK$            // case 1:  $x$ 's sibling  $w$  is red.
6               $x.p.color = RED$            // case 1
7              LEFT-ROTATE( $T, x.p$ )       // case 1
8               $w = x.p.right$            // case 1
9          if  $w.left.color == BLACK$  and  $w.right.color == BLACK$ 
10              $w.color = RED$            // case 2:  $x$ 's sibling  $w$  is black, and ...
11              $x = x.p$                  // case 2: ... both of  $w$ 's children ...
                                     // case 2: ... are black.
```


Continued on next page.

RB-DELETE-FIXUP

Pseudocode (Part 2)

```
12         else if  $w.right.color == \text{BLACK}$ 
13              $w.left.color = \text{BLACK}$  // case 3:  $x$ 's sibling  $w$  is black, ...
14              $w.color = \text{RED}$  // case 3: ...  $w$ 's left child is red, ...
15              $\text{RIGHT-ROTATE}(T, w)$  // case 3: ... and  $w$ 's right child ...
16              $w = x.p.right$  // case 3: ... is black.
17              $w.color = x.p.color$  // case 4:  $x$ 's sibling  $w$  is black, and ...
18              $x.p.color = \text{BLACK}$  // case 4: ...  $w$ 's right child is red.
19              $w.right.color = \text{BLACK}$  // case 4
20              $\text{LEFT-ROTATE}(T, x.p)$  // case 4
21              $x = T.root$  // case 4
22         else (same as then clause with "right" and "left" exchanged)
23      $x.color = \text{BLACK}$ 
```

Case 1


 x 's Sibling w Is Red

Since w must have black children, we can switch the colors of w and $x.p$ and then perform a left-rotation on $x.p$ without violating any of the red-black properties. The new sibling of x , which is one of w 's children prior to the rotation, is now black, and thus we have converted case 1 into case 2, 3, or 4.

Case 2

x 's Sibling w Is Black, and Both of w 's Children Are Black

Both light grey nodes labeled B in the diagrams below have the same color attribute c , which may be either RED or BLACK.




When we arrive at case 2, the black-height of w equals the black-height of x plus 1. We reduce the black-height of w by making it red. Afterwards, x moves up to its former parent.

If we arrive at case 2 through case 1, the new node x is red. Hence, the **while** loop in RB-DELETE-FIXUP terminates and x changes its color to black.

Case 3

x 's Sibling w Is Black, w 's Left Child Is Red, and w 's Right Child Is Black




We can switch the colors of w and its left child $w.left$ and then perform a right rotation on w without violating any of the red-black properties. The new sibling w of x is now a black node with a red right child. Therefore, we have transformed case 3 into case 4.

Case 4

x's Sibling w Is Black, and w 's Right Child Is Red

The color attributes c and c' in the diagram below may be either RED or BLACK.



When we arrive at case 4, the black-height of w equals the black-height of x plus 1. By making some color changes and performing a left rotation on $x.p$, we balance the black-heights and also satisfy all other red-black properties.

Setting x to be the root causes the **while** loop to terminate.

Running Time of RB-DELETE

$O(\log n)$

Since the height of a red-black tree of n nodes is $O(\log n)$, the total cost of the procedure without the call to RB-DELETE-FIXUP takes $O(\log n)$ time.


Within RB-DELETE-FIXUP, each of cases 1, 3, and 4 lead to termination after performing a constant number of color changes and at most three rotations. Case 2 is the only case in which the **while** loop can be repeated, and then the pointer x moves up the tree at most $O(\log n)$ times, performing no rotations.

Therefore, the procedure RB-DELETE-FIXUP takes $O(\log n)$ time and performs at most three rotations, and the overall time for RB-DELETE is also $O(\log n)$.

Exercise

Deleting Nodes from a Red-Black Tree

Consider the initial red-black tree shown below. Show the red-black trees that result from the successive deletion of the keys in the order 8, 12, 19, 31, 38, 41.



Solution to Exercise

<https://github.com/gzc/CLRS/blob/master/C13-Red-Black-Trees/13.4.md>

